

R INTRODUCTION

Tobias Niedermaier

INTRO

WHY R?

Why R?

- R is **open source**
- All techniques for data analyses
- State-of-the-art graphics capabilities
- A platform for programming new statistical methods or analysis pipelines (in form of R-packages)

PROGRAMMING (IN GENERAL)

“Good programmers are made, not born.” (Gerald M. Weinberg - The Psychology of Computer Programming)

- consequence I
train...
- consequence II
train...
- consequence III
train more

Hands-on is important. Understanding is less than 30%

R AND R STUDIO

Required tools for the course:

- Programming language R
 - designed to make fast prototyping for statistical analysis
 - interpreted language
- RStudio (optional, but recommended)
 - IDE tailored for R
 - Integrates a lot more (e.g. python, c++, etc.)
- Alternatively, you may use any code editor that supports R and copy-paste your codes into the command line (Notepad++, Emacs etc.)

R PACKAGES

- R comes with many useful packages by default
- However, the strength lies in the huge collection of external packages
- Most popular and default: CRAN
- Install new packages in R using either
 - using a command:
 - `install.packages("<package-name>")`
(e.g. `install.packages("mvtnorm")`)
 - RStudio
 - using built-in tools from the IDE

BASIC OPERATIONS

ADDITION, SUBTRACTION, ETC

```
1 1+2
```

```
[1] 3
```

```
1 1-2
```

```
[1] -1
```

```
1 1*2
```

```
[1] 2
```

```
1 1/2
```

```
[1] 0.5
```

```
1 1^2
```

```
[1] 1
```

```
1 1**2
```

```
[1] 1
```

Note

What will happen?

```
1 1/0
```


SPECIAL SYMBOLS FUNCTIONS

Special symbols

```
1 pi
2 Inf
```

Mathematical functions

```
1 exp(1)
```

```
[1] 2.718282
```

```
1 log(1)
```

```
[1] 0
```

Special cases:

- **NaN** is a data type that indicates an invalid number.

```
1 log(-1)
```

```
[1] NaN
```

```
1 NaN + 1
```

```
[1] NaN
```

- **NA** is a missing value.

```
1 NA + 1
```

```
[1] NA
```

- **NULL** means literally empty/nothing

ASSIGNING OBJECTS

Assignment is done using `<-`

```
1 x <- 1  
2 y <- 2  
3 x + y
```

```
[1] 3
```

Alternatively, use `=`

```
1 x = sqrt(2)  
2 y = sqrt(2)  
3 x * y
```

```
[1] 2
```

Or (almost never used):

```
1 5 -> x
```

Look at environment pane in R Studio, what can you see?

Also try:

```
1 ls()
```

```
[1] "x" "y"
```

```
1 ls.str()
```

```
x :   num 5
```

```
y :   num 1.41
```

NAMING OBJECTS

- Objects in **R** have to start with a letter

Case sensitive

```
1 a <- 2
2 A <- 1
3 a-A
```

```
[1] 1
```

Overwrite variables with old ones

```
1 a <- a + 1 #There is NO warning if you overwrite an existing object!
```

Combination of words

```
1 variable_name <- 1
2 variable.name <- 1
3 variableName <- 1
```

COMMENTS

Sometimes it is useful, to comment code. Use a `#` to comment

Standard:

```
1 1+1
```

```
[1] 2
```

Comment a line (no output):

```
1 # 1+1
```

Comment after an expression (only `1+1` gets evaluated):

```
1 1+1 # +1
```

```
[1] 2
```

FUNCTION CALLING

So far we used expressions like `f(...)`. This is a **function**. E.g.

```
1 exp(2)
```

We call the function `exp` with a value of 2. Or the (natural) logarithm:

```
1 log(exp(1))
```

```
[1] 1
```

We can specify the base as a *second argument*:

```
1 log(2, 2)
```

```
[1] 1
```

Note

What will happen?

```
1 Log(Exp(1))
```

GET DOCUMENTATION

Access the documentation using

- `<F1>`
- type `?function_name`
- use RStudio functionality

E.g. documentation for `log()` reveals that we calculate the natural logarithm.

```
1 ?log
2 log(x, base = exp(1))
```


FUNCTION CALLING CONT'D

You can ignore the argument name, when placements are clear. - We have done that for `exp` and `log`

Hence, this here

```
1 log(2, 2)
```

means, that we actually call

```
1 log(x=2, base=2)
```

If you specify the argument, order does not matter.

Example:

```
1 log(base=3, x=2)
2 log(3, 2)
```

Note

What will happen?

```
1 log <- 1
2 log(log)
```

BASIC (PRIMITIVE) DATA TYPES

numeric

A (floating point) number. We used this so far (default).

1.0, 1.34, -33, pi

logical

A binary data type.

TRUE, FALSE, T, F

Note

What will happen?

```
1 sum(c(TRUE, FALSE, FALSE, TRUE, FALSE))
2 sum(!c(TRUE, FALSE, FALSE, TRUE, FALSE))
```

integer

Can be specified using an “L”.

1L, 100L, -99L

character

Represents letters OR sentences.

'a', "abc", "May the force be with you"

EXERCISES 1 TASK 1

VECTORS

VECTORS

You can *combine* single values to a *vector*.

```
1 a <- c(1,2,3,4)
2 a
```

```
[1] 1 2 3 4
```

```
1 b <- c(TRUE, FALSE, TRUE)
2 b
```

```
[1] TRUE FALSE TRUE
```

```
1 c <- c("a", 'ab', "ab c")
2 c
```

```
[1] "a"      "ab"     "ab c"
```

Many operations in R are *vectorized*

```
1 a + a
```

```
[1] 2 4 6 8
```

```
1 a * a
```

```
[1] 1 4 9 16
```

```
1 exp(a)
```

```
[1] 2.718282 7.389056 20.085537 54.598150
```

```
1 -a
```

```
[1] -1 -2 -3 -4
```

Note

What will happen?

```
1 c("1",2,3)
```

VECTORS CONT'D

- NA or NaNs can be part of a vector

```
1 a <- c(1,2,NA,4)
2 a + 1
```

```
[1] 2 3 NA 5
```

```
1 b <- c(1, -1, Inf)
2 log(b)
```

```
[1] 0 NaN Inf
```


AUTOMATIC RECYCLING

```
1 a <- c(1,2,3,4)
2 a + 1
```

```
[1] 2 3 4 5
```

```
1 b <- c(2,2)
2 a + b
```

```
[1] 3 4 5 6
```

Warning

Note the behavior for vectors with different length! Example:

```
1 a <- c(1,2,3)
2 b <- c(1,2)
3 a + b
```

Warning in a + b: Länge des längeren Objektes
ist kein Vielfaches der Länge des kürzeren Objektes

```
[1] 2 4 4
```

VECTOR CREATION

There are a lot of convenience functions to create vectors.

```
1 c(1,2,3,4)
```

```
[1] 1 2 3 4
```

```
1 1:4
```

```
[1] 1 2 3 4
```

```
1 seq(4)
```

```
[1] 1 2 3 4
```

More complex ones:

```
1 4:-3
```

```
[1] 4 3 2 1 0 -1 -2 -3
```

```
1 seq(-10, 10, by = 2)
```

```
[1] -10 -8 -6 -4 -2 0 2 4 6 8 10
```

```
1 seq(-10, 10, length.out = 10) # vector of length 10
```

```
[1] -10.000000 -7.777778 -5.555556 -3.333333 -1.111111 1.111111
```

```
[7] 3.333333 5.555556 7.777778 10.000000
```

SELECT ELEMENTS OF A VECTOR

Access elements of a vector using positional numbers within [...]:

```
1 x <- c(2,4,2,5)
2 x[1]
```

```
[1] 2
```

Multiple elements

```
1 selection <- c(1,4)
2 x[selection]
```

```
[1] 2 5
```

```
1 x[c(1,4)]
```

```
[1] 2 5
```

Negative values will be excluded

```
1 x[-c(1,3)]
```

```
[1] 4 5
```

Note

What will happen?

```
1 x[1:5]
2 x[-(5:10)]
3 x[0]
```

LOGICAL VALUES FOR COMPARISON

Recall the very most basic data type `logical`, i.e. `TRUE` and `FALSE`.

- We can create such an object by comparison:

```
1 1 == 2 # lhs equal rhs?
```

```
[1] FALSE
```

```
1 1 != 2 # lhs unequal rhs?
```

```
[1] TRUE
```

```
1 1 > 2 # lhs larger rhs?
```

```
[1] FALSE
```

```
1 1 >= 2 # lhs larger or equal rhs?
```

```
[1] FALSE
```

```
1 1 < 2 # lhs less than rhs?
```

```
[1] TRUE
```

```
1 1 <= 2 # lhs less or equal than rhs?
```

```
[1] TRUE
```

Swap value:

```
1 !TRUE
```

```
[1] FALSE
```

```
1 !FALSE
```

```
[1] TRUE
```

Note

What will happen?

```
1 1 == "1"  
2 1 != NaN  
3 NA == NA # we will learn the solution in a few slides
```

FILTER ELEMENTS OF A VECTOR

Comparison operators are vectorized:

```
1 c(T,F,T) == c(F,F,T) # element-wise comparison
```

```
[1] FALSE TRUE TRUE
```

Except if you use `identical()` which is `FALSE` if there is any mismatch:

```
1 identical(c(T,F,T), c(F,F,T))
```

```
[1] FALSE
```

Check condition on a numeric vector

```
1 x <- c(2,4,2,5)
2 position_two <- x == 2 # logical vector showing, where the condition holds
3 position_two
```

```
[1] TRUE FALSE TRUE FALSE
```

Use logical values to filter a vector.

```
1 x[position_two]
```

```
[1] 2 2
```

```
1 # or directly
2 x[x == 2]
```

```
[1] 2 2
```

Filter for values less than 3

```
1 x[x < 3]
```

```
[1] 2 2
```


COMBINE FILTERS WITH & AND |

Combination operations...

```
1 TRUE & TRUE
```

```
[1] TRUE
```

```
1 FALSE & TRUE
```

```
[1] FALSE
```

```
1 TRUE | TRUE
```

```
[1] TRUE
```

```
1 FALSE | TRUE
```

```
[1] TRUE
```

...or vectorized

```
1 x <- c(T,F,T,F)
2 y <- c(T,T,F,F)
3 x & y
```

```
[1] TRUE FALSE FALSE FALSE
```

```
1 x | y
```

```
[1] TRUE TRUE TRUE FALSE
```

Use this to filter a vector for multiple conditions

```
1 x[(x < 5) & (x > 2)]
```

```
logical(0)
```

ASSIGN NEW VALUES IN A VECTOR

We can assign new values to a vector using a combination of selection and assignment

```
1 x <- 1:5  
2 x[1] <- 2  
3 x
```

```
[1] 2 2 3 4 5
```

```
1 x[x > 3] <- -99  
2 x
```

```
[1] 2 2 3 -99 -99
```

```
1 x[-1] <- 100  
2 x
```

```
[1] 2 100 100 100 100
```

Note

What will happen?

```
1 x[100] <- 1
```

VECTOR OPERATIONS

```
1 x <- c(1,1,2,3)
2 length(x)
```

```
[1] 4
```

```
1 append(x, c(1,2,3))
```

```
[1] 1 1 2 3 1 2 3
```

```
1 rev(x)
```

```
[1] 3 2 1 1
```

```
1 sort(x)
```

```
[1] 1 1 2 3
```

```
1 unique(x)
```

```
[1] 1 2 3
```

```
1 sum(x)
```

```
[1] 7
```

EXERCISES 1 TASK 2

COMPLEX STRUCTURES

FACTORS

Consider a vector, that represents a categorical variable. Let's say colors.

```
1 colors <- c("blue", "red", "blue", "red", "green", "black", "green", "white")
2 colors
```

```
[1] "blue" "red" "blue" "red" "green" "black" "green" "white"
```

We cast `colors` into a factor now:

```
1 colors <- as.factor(colors)
2 colors
```

```
[1] blue red blue red green black green white
Levels: black blue green red white
```

```
1 levels(colors)
```

```
[1] "black" "blue" "green" "red" "white"
```

```
1 as.numeric(colors)
```

```
[1] 2 4 2 4 3 1 3 5
```

```
1 class(colors)
```

```
[1] "factor"
```

```
1 typeof(colors)
```

```
[1] "integer"
```

COMPLEX DATA STRUCTURES

from *Ceballos and Cardiel, (2013). Data structure – First Steps in R. Retrieved 25-11-2018 from http://venus.ifca.unican.es/Rintro/_images/dataStructuresNew.png*

Use `str(...)` to inspect the structure of complex data types!

VECTOR, MATRIX, ARRAY

We already got vectors. Lets combine them:

```
1 x <- 1:4
2 (x_rbind <- cbind(x,x)) # 4 rows, 2 columns
```

```
      x x
[1,] 1 1
[2,] 2 2
[3,] 3 3
[4,] 4 4
```

```
1 (x_cbind <- rbind(x,x)) # 2 rows, 4 columns
```

```
  [,1] [,2] [,3] [,4]
x     1     2     3     4
x     1     2     3     4
```

```
1 dim(x_rbind)
```

```
[1] 4 2
```

```
1 dim(x_cbind)
```

```
[1] 2 4
```

```
1 nrow(x_rbind)
```

```
[1] 4
```

```
1 ncol(x_rbind)
```

```
[1] 2
```

VECTOR, MATRIX, ARRAY CONT'D

We can define a matrix using the `matrix` function:

```
1 matrix(1:6, nrow = 3, ncol = 2)
```

```
      [,1] [,2]  
[1,]     1     4  
[2,]     2     5  
[3,]     3     6
```

```
1 matrix(1:6, nrow = 3, ncol = 2, byrow = T)
```

```
      [,1] [,2]  
[1,]     1     2  
[2,]     3     4  
[3,]     5     6
```

Arrays as a generalization with multiple dimensions

```
1 array(1:12, dim = c(3,2,2))
```

```
, , 1
```

```
      [,1] [,2]  
[1,]     1     4  
[2,]     2     5  
[3,]     3     6
```

```
, , 2
```

	[,1]	[,2]
[1,]	7	10
[2,]	8	11
[3,]	9	12

This is also sometimes called a *tensor*.

SELECT/FILTER ELEMENTS ON ARRAYS

As vectors, we can select and filter. Seperate dimensions with a `,`,
i.e. `[... , ...]`

```
1 (m <- matrix(1:6, nrow = 3, ncol = 2))
```

```
      [,1] [,2]  
[1,]    1    4  
[2,]    2    5  
[3,]    3    6
```

```
1 m[2,2]
```

```
[1] 5
```

```
1 m[nrow(m), ncol(m) ]
```

```
[1] 6
```

Defining no entry will return the full dimension:

```
1 m[2,]
```

```
[1] 2 5
```

```
1 m[,1]
```

```
[1] 1 2 3
```

Note

What will happen?

```
1 m[1,,2]  
2 m[10]
```

LIST

A list is a collection of elements. These elements could be any object.

```
1 (l <- list(1, "2", 1:3, list(m)))
```

```
[[1]]
```

```
[1] 1
```

```
[[2]]
```

```
[1] "2"
```

```
[[3]]
```

```
[1] 1 2 3
```

```
[[4]]
```

```
[[4]][[1]]
```

```
 [,1] [,2]
```

```
[1,]    1    4
```

```
[2,]    2    5
```

```
[3,]    3    6
```

Access elements of a list with `[[...]]`.

```
1 l[[2]]
```

```
[1] "2"
```

A sub-list can be accessed with `[[...]]`.

```
1 1[1:3]
```

```
[[1]]
```

```
[1] 1
```

```
[[2]]
```

```
[1] "2"
```

```
[[3]]
```

```
[1] 1 2 3
```

LIST CONT'D

You can define names for lists:

```
1 l <- list(slot1 = 1:3, slot2 = c("a", "b"), slot3 = 1)
2 names(l)
```

```
[1] "slot1" "slot2" "slot3"
```

Access list elements using the name and a `$`:

```
1 l$slot3 # return the original list l before overwriting it
```

```
[[1]]
```

```
[1] 1
```

```
[[2]]
```

```
[1] "2"
```

```
[[3]]
```

```
[1] 1 2 3
```

```
[[4]]
```

```
[[4]][[1]]
```

```
  [,1] [,2]
```

```
[1,]   1   4
```

```
[2,]   2   5
```

```
[3,]   3   6
```


Delete elements by assigning a **NULL** to a slot

```
1 1[2:3] <- NULL  
2 1
```

```
$slot1
```

```
[1] 1 2 3
```

DATA FRAME

A data frame is basically a list, where each element is a vector of the same length. However, it implements function to handle it as a matrix.

Let's define a data set representing cars:

```
1 col <- as.factor(c("blue", "red", "blue", "red", "green", "black", "green", "white"))
2 pri <- c(10, 20, 9, 50, 0.4, 15, 160, 60) * 1000
3 is_el <- c(F,F,F,T,F,T,F,T)
4
5 car_ds <- data.frame(color = col, price = pri, is_electric = is_el)
6 car_ds
```

	color	price	is_electric
1	blue	10000	FALSE
2	red	20000	FALSE
3	blue	9000	FALSE
4	red	50000	TRUE
5	green	400	FALSE
6	black	15000	TRUE
7	green	160000	FALSE
8	white	60000	TRUE

```
1 str(car_ds)
```

```
'data.frame':  8 obs. of  3 variables:
 $ color      : Factor w/ 5 levels "black","blue",...: 2 4 2 4 3 1 3 5
 $ price      : num  10000 20000 9000 50000 400 15000 160000 60000
 $ is_electric: logi  FALSE FALSE FALSE TRUE FALSE TRUE
```

DATA FRAME CONT'D

We can work on a data set as we work with a matrix

```
1 # All rows with red cars
2 car_ds[car_ds$color == "red", ]
```

```
   color price is_electric
2   red 20000         FALSE
4   red 50000          TRUE
```

```
1 # price of all black cars
2 car_ds[car_ds$color == "black", "price"]
```

```
[1] 15000
```

```
1 # set a new price for the last car in the ds
2 car_ds[8, 2] <- 600
3 car_ds
```

```
   color price is_electric
1  blue  10000         FALSE
2   red  20000         FALSE
3  blue   9000         FALSE
4   red  50000          TRUE
5 green   400         FALSE
6 black 15000          TRUE
7 green 160000         FALSE
8 white   600          TRUE
```

MORE ON DATA STRUCTURES

- A data frame behaves like a matrix.
- However, keep in mind that it is actually a list. We can easily prove that:

```
1 is.list(car_ds)
```

```
[1] TRUE
```

Use `str(...)` to check the data structure of *any* object:

```
1 str(car_ds)
```

```
'data.frame':  8 obs. of  3 variables:
 $ color      : Factor w/ 5 levels "black","blue",...: 2 4 2 4 3 1 3 5
 $ price      : num  10000 20000 9000 50000 400 15000 160000 600
 $ is_electric: logi  FALSE FALSE FALSE TRUE FALSE TRUE ...
```

```
1 m <- matrix(1:4, ncol = 2)
2 str(m)
```

```
int [1:2, 1:2] 1 2 3 4
```

You will frequently need conditional subsetting, combining several conditions:

```
1 mtcars[mtcars$mpg > 15 & mtcars$wt < 3 & mtcars$hp <= 66,]
```

	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
Fiat 128	32.4	4	78.7	66	4.08	2.200	19.47	1	1	4	1
Honda Civic	30.4	4	75.7	52	4.93	1.615	18.52	1	1	4	2
Toyota Corolla	33.9	4	71.1	65	4.22	1.835	19.90	1	1	4	1
Fiat X1-9	27.3	4	79.0	66	4.08	1.935	18.90	1	1	4	1

LOAD DATA

We can load a data set from a package using `data(...)`.

```
1 data("iris", package = "datasets") # look in the environment variables
```

We can load data from files. Use `read.table(...)`, or wrapper functions with reasonable default values. E.g. We can read a file directly from the web:

```
1 d <- read.csv("https://raw.githubusercontent.com/vincentarelbundock/Rdatasets/master/csv/datasets/m
2 head(d) # show the first few lines of a data set
```

	rownames	mpg	cyl	disp	hp	drat	wt	qsec	vs	am	gear	carb
1	Mazda RX4	21.0	6	160	110	3.90	2.620	16.46	0	1	4	4
2	Mazda RX4 Wag	21.0	6	160	110	3.90	2.875	17.02	0	1	4	4
3	Datsun 710	22.8	4	108	93	3.85	2.320	18.61	1	1	4	1
4	Hornet 4 Drive	21.4	6	258	110	3.08	3.215	19.44	1	0	3	1
5	Hornet Sportabout	18.7	8	360	175	3.15	3.440	17.02	0	0	3	2
6	Valiant	18.1	6	225	105	2.76	3.460	20.22	1	0	3	1

Note, that we can also use this to read a data set from a local directory! To do that we have to specify either the full path or define the path from the **working directory**. Use `getwd(...)` and `setwd(...)` to get or set the current working directory. See next slide for an example.

SAVE DATA SETS

Consider a data set, you have worked with. You can save it using write functions.

```
1 write.csv(car_ds, file = "example_data.csv") # we save our data set in the current working directory
```

We can again read the data as a new object:

```
1 d_loaded <- read.csv("example_data.csv")
2
3 all.equal(car_ds, d_loaded) # test whether 2 (more complex) R object are the same
```

```
[1] "Names: 3 string mismatches"
[2] "Length mismatch: comparison on first 3 components"
[3] "Component 1: 'current' is not a factor"
[4] "Component 2: Modes: numeric, character"
[5] "Component 2: target is numeric, current is character"
[6] "Component 3: Modes: logical, numeric"
[7] "Component 3: target is logical, current is numeric"
```

We can read other files as well. E.g. excel, SPSS, SAS, etc.

There are a lot of packages to do that.

I use the function `load(...)` from the `rio` package that tries to unify a lot of different formats)

SAVE AND LOAD R OBJECTS

So far, we only worked with data frames for read and write operations. We can save general R objects using `save(...)` and `load(...)` using the `.RData` format.

```
1 a_list <- list(a = 42, data = iris, comment = "whatever")
2
3 save(a_list, file = "example_object.RData")
4
5 load("example_object.RData")
```


EXERCISES 1 TASKS 3